

Class 7

Sabrina Koldinger (A16368238)


Class 7 Clustering methods

The broad goal here is to find groupings in your input data.

kmeans

First, make some data to cluster.

```
x=rnorm(1000)  
hist(x)
```




MAke a vector of length 60 with 30 points centered at -3 and 30 points centered at 3.

```
tmp=c(rnorm(30, mean = -3), rnorm(30, mean = 3))
tmp
```

```
[1] -1.993215 -2.545352 -3.672325 -2.906536 -3.095059 -3.517621 -3.777858
[8] -1.832730 -2.063347 -3.761910 -3.567369 -2.284752 -3.692237 -2.978851
[15] -1.638371 -2.640830 -3.262984 -2.595547 -3.577998 -3.182104 -3.286236
[22] -3.143687 -3.435739 -3.203561 -3.735788 -2.044038 -3.778846 -3.589206
[29] -2.874677 -3.175576  3.283563  4.212834  2.485628  1.917040  3.955702
[36]  5.343970  4.662216  2.426890  2.699276  2.596221  3.428800  4.200494
[43]  4.457996  3.486060  2.774215  2.552768  3.815276  2.958822  3.508970
[50]  4.031265  2.935926  2.504427  3.276542  3.280792  2.344021  2.796407
[57]  3.413238  1.747618  2.768272  3.554944
```

I will now make a small x and y data set with 2 groups of points.

```
x= cbind(x=tmp, y=rev(tmp))
plot(x)
```



```
k=kmeans(x,centers=2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.028478	3.247340
2	3.247340	-3.028478

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 32.15525 32.15525
(between_SS / total_SS = 94.8 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Question: from your result object `k` how many point are in each cluster?

k\$size

[1] 30 30

Question: What “component” of your result object details the cluster membership?

```
k$cluster
```

[illegible]


Question: what about cluster center?

```
k$centers
```

	x	y
1	-3.028478	3.247340
2	3.247340	-3.028478

Plot clusters

```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch=15, cex=2)
```



We can cluster into 4 groups

```
k4=kmeans(x,centers=4)
k4
```

K-means clustering with 4 clusters of sizes 9, 14, 30, 7

Cluster means:

	x	y
1	3.145453	-2.216575
2	2.741933	-3.455099
3	-3.028478	3.247340
4	4.389150	-3.219113

Clustering vector:

```
[1] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 2 4 2 2 1 4 4 2  
[39] 2 2 2 4 4 2 1 1 4 2 1 4 2 1 1 2 2 2 1 2 1 1
```


Within cluster sum of squares by cluster:

```
[1] 3.441073 4.585824 32.155249 2.597549  
(between_SS / total_SS = 96.6 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(x, col=k4$cluster)
```



A big limitation of kmeans is that it does what you ask even if the clusters are silly.

hclust

The main base R function for Hierarchical clustering is `hclust`. Unlike `kmeans`, you can not just pass it your data as an input. You need to calculate a distance matrix.

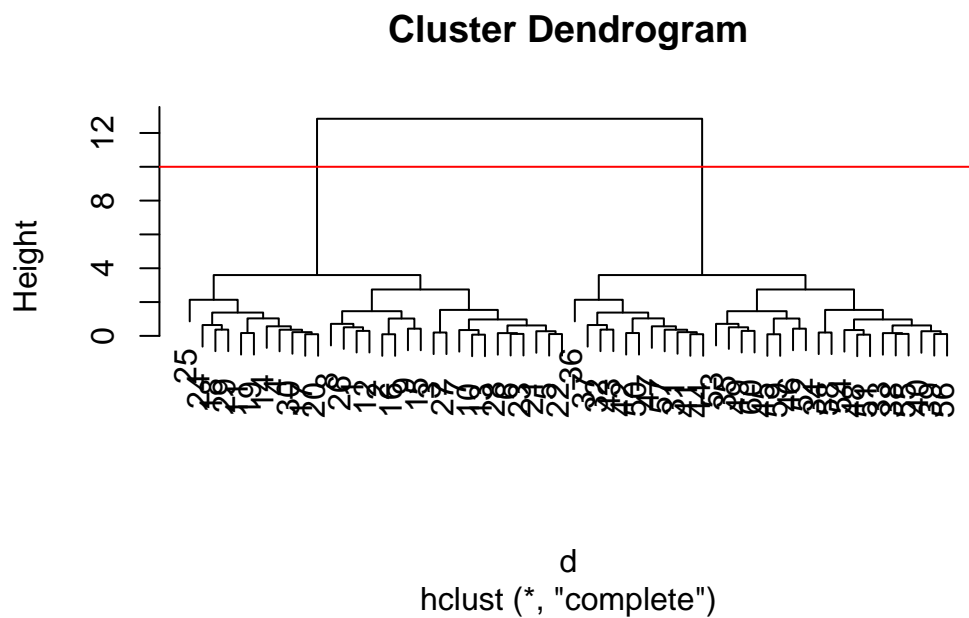
```
d=dist(x)
hc=hclust(d)
hc
```

Call:
`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

Use `plot()`

```
plot(hc)
abline(h=10, col="red")
```




To make the “cut” and get our cluster membership vector we can use the `cutree()` function. Make a plot of our data colored by hclust results.

```
groups=cutree(hc, h=10)
groups
```

[illegible]

```
plot(x, col=groups)
```



PCA- Principal Component Analysis

Here we will do Principal Component Analysis (PCA) on some food data from the UK.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

```
#dont do this
#rownames(x)=x[,1]
#x=x[,-1]
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```


17 rows and 4 columns

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

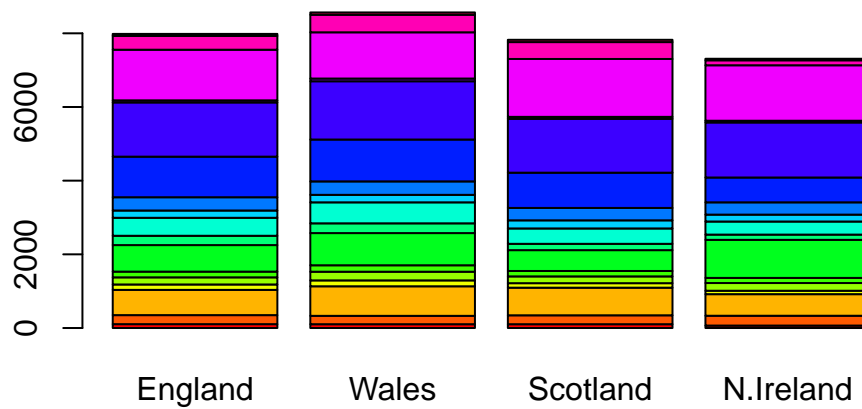
It is more effective to use the `row.names=1` since it will not end up removing columns of data if run multiple times.

Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



```
pairs(x, col=rainbow(10), pch=16)
```



PCA to the rescue

The main “base” R function for PCA is called `prcomp()`.

```
pca=prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

How much variance is captured in 2 PCs. 96.5%

To make our main “PC score plot” or PC1 vs PC2 plot”, “PC plot”, or ” ordination plot”.

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class


```
[1] "prcomp"
```

We are after the `pca$x` result component to make our main PCA plot.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
mycols=c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab="PC1(67.4%)", ylab="PC2(29%)")
```



Another important result from PCA is how the original variables (in this case foods) contribute to the PCs. This is contained in the `pca$rotation` object- people often call this the “loadings” or “contributions” to the PCs.

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484

Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001890737

We can make a plot along PC1

```
library(ggplot2)
contrib= as.data.frame(pca$rotation)
ggplot(contrib)+ aes(PC1, rownames(contrib))+geom_col()
```

